

Time Series

Packages

Uses my package `mkac` which is on Github. Install with:

```
library(devtools)
install_github("nxskok/mkac")
```

Plus these. You might need to install some of them first:

```
library(ggfortify)
library(forecast)
library(tidyverse)
library(mkac)
```


Time trends

- Assess existence or nature of time trends with:
 - correlation
 - regression ideas.
 - (later) time series analysis

World mean temperatures

Global mean temperature every year since 1880:

```
temp=read_csv("temperature.csv")  
ggplot(temp, aes(x=year, y=temperature)) +  
  geom_point() + geom_smooth()
```



Examining trend

- Temperatures increasing on average over time, but pattern very irregular.
- Find (Pearson) correlation with time, and test for significance:

```
with(temp, cor.test(temperature,year))
```

```
##  
## Pearson's product-moment correlation  
##  
## data: temperature and year  
## t = 19.996, df = 129, p-value < 2.2e-16  
## alternative hypothesis: true correlation is not equal to 0  
## 95 percent confidence interval:  
## 0.8203548 0.9059362  
## sample estimates:  
## cor  
## 0.8695276
```

Comments

- Correlation, 0.8695, significantly different from zero.
- CI shows how far from zero it is.

Tests for *linear* trend with *normal* data.

Kendall correlation

Alternative, Kendall (rank) correlation, which just tests for monotone trend (anything upward, anything downward) and is resistant to outliers:

```
with(temp, cor.test(temperature,year,method="kendall"))
```

```
##  
## Kendall's rank correlation tau  
##  
## data: temperature and year  
## z = 11.776, p-value < 2.2e-16  
## alternative hypothesis: true tau is not equal to 0  
## sample estimates:  
## tau  
## 0.6992574
```

Kendall correlation usually closer to 0 for same data, but here P-values comparable. Trend again strongly significant.

Mann-Kendall

- Another way is via **Mann-Kendall**: Kendall correlation with time.
- Use my package mkac:

```
kendall_Z_adjusted(temp$temperature)
```

```
## $z
## [1] 11.77267
##
## $z_star
## [1] 4.475666
##
## $ratio
## [1] 6.918858
##
## $P_value
## [1] 0
##
## $P_value_adj
## [1] 7.617357e-06
```


Comments

- Standard Mann-Kendall assumes observations *independent*.
- Observations close together in time often *correlated* with each other.
- Correlation of time series “with itself” called **autocorrelation**.
- Adjusted P-value above is correction for autocorrelation.

Examining rate of change

- Having seen that there *is* a change, question is “how fast is it?”
- Examine slopes:
 - regular regression slope, if you believe straight-line regression
 - Theil-Sen slope: resistant to outliers, based on medians

Ordinary regression against time

```
lm(temperature~year, data=temp) %>% tidy() -> temp.tidy  
temp.tidy
```

term	estimate	std.error	statistic	p.value
(Intercept)	2.5794197	0.5703984	4.522137	1.37e-05
year	0.0058631	0.0002932	19.996448	0.00e+00

- Slope about 0.006 degrees per year
- about this many degrees over course of data):

```
temp.tidy %>% pluck("estimate", 2)*130
```

```
## [1] 0.7622068
```

Theil-Sen slope

also from mkac:

```
theil_sen_slope(temp$temperature)
```

```
## [1] 0.005675676
```


Conclusions

- Slopes:
 - Linear regression: 0.005863
 - Theil-Sen slope: 0.005676
 - Very close.
- Correlations:
 - Pearson 0.8675
 - Kendall 0.6993
 - Kendall correlation smaller, but P-value equally significant (often the case)

Constant rate of change?

Slope assumes that the rate of change is same over all years, but trend seemed to be accelerating:

```
ggplot(temp, aes(x=year, y=temperature)) +  
  geom_point() + geom_smooth()
```



Pre-1970 and post-1970:

```
temp %>%
  mutate(time_period=
    ifelse(year<=1970, "pre-1970", "post-1970")) %>%
  nest(-time_period) %>%
  mutate(theil_sen=map_dbl(
    data, ~theil_sen_slope(.$temperature)))
```

```
## Warning: All elements of `...` must be named.
```

```
## Did you want `data = c(X1, Year, temperature, year)`?
```

time_pe-
riodata

pre1970	1.00	2.00	3.00	4.00	5.00	6.00	7.00	8.00	9.00	10.00	11.00	12.00	13.00	14.00	15.00	16.00	17.00	18.00	19.00	20.00	21.00	22.00	23.00	24.00	25.00	26.00	27.00	28.00	29.00	30.00	31.00	32.00	33.00	34.00	35.00	36.00	37.00	38.00	39.00	40.00	41.00	42.00	43.00	44.00	45.00	46.00	47.00	48.00	49.00	50.00	51.00	0.00
---------	------	------	------	------	------	------	------	------	------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	------

Actual time series: the Kings of England

- Age at death of Kings and Queens of England since William the Conqueror (1066):

```
kings=read_table("kings.txt", col_names=F)
```

```
##  
## -- Column specification -----  
## cols(  
##   X1 = col_double()  
## )
```

Data in one long column X1, so kings is data frame with one column.

Turn into ts time series object


```
kings.ts=ts(kings)
kings.ts
```

```
## Time Series:
## Start = 1
## End = 42
## Frequency = 1
##      X1
## [1,] 60
## [2,] 43
## [3,] 67
## [4,] 50
## [5,] 56
## [6,] 42
## [7,] 50
## [8,] 65
## [9,] 60
## [10,] 50
```

Plotting a time series

autoplot from ggfortify gives time plot:

```
autoplot(kings.ts)
```



Comments

- “Time” here is order of monarch from William the Conqueror (1st) to George VI (last).
- Looks to be slightly increasing trend of age-at-death
- but lots of irregularity.

Stationarity

A time series is **stationary** if:

- mean is constant over time
- variability constant over time and not changing with mean.

Kings time series seems to have:

- non-constant mean
- but constant variability
- not stationary.

Getting it stationary

- Usual fix for non-stationarity is *differencing*: get new series from original one's values: 2nd - 1st, 3rd - 2nd etc.


In R, diff:

```
kings.diff.ts=diff(kings.ts)
```

Did differencing fix stationarity?

Looks stationary now:

```
autoplot(kings.diff.ts)
```



Births per month in New York City

from January 1946 to December 1959:

```
ny=read_table("nybirths.txt",col_names=F)  
ny
```

X1
26.663
23.598
26.931
24.740
25.806
24.364
24.477
23.901
23.175
23.227
21.672
21.870

As a time series

```
ny.ts=ts(ny,freq=12,start=c(1946,1))  
ny.ts
```

##		Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
##	1946	26.663	23.598	26.931	24.740	25.806	24.364	24.477	23.901	23.175	23.227	21.672	21.870
##	1947	21.439	21.089	23.709	21.669	21.752	20.761	23.479	23.824	23.105	23.110	21.759	22.073
##	1948	21.937	20.035	23.590	21.672	22.222	22.123	23.950	23.504	22.238	23.142	21.059	21.573
##	1949	21.548	20.000	22.424	20.615	21.761	22.874	24.104	23.748	23.262	22.907	21.519	22.025
##	1950	22.604	20.894	24.677	23.673	25.320	23.583	24.671	24.454	24.122	24.252	22.084	22.991
##	1951	23.287	23.049	25.076	24.037	24.430	24.667	26.451	25.618	25.014	25.110	22.964	23.981
##	1952	23.798	22.270	24.775	22.646	23.988	24.737	26.276	25.816	25.210	25.199	23.162	24.707
##	1953	24.364	22.644	25.565	24.062	25.431	24.635	27.009	26.606	26.268	26.462	25.246	25.180
##	1954	24.657	23.304	26.982	26.199	27.210	26.122	26.706	26.878	26.152	26.379	24.712	25.688
##	1955	24.990	24.239	26.721	23.475	24.767	26.219	28.361	28.599	27.914	27.784	25.693	26.881
##	1956	26.217	24.218	27.914	26.975	28.527	27.139	28.982	28.169	28.056	29.136	26.291	26.987
##	1957	26.589	24.848	27.543	26.896	28.878	27.390	28.065	28.141	29.048	28.484	26.634	27.735
##	1958	27.132	24.924	28.963	26.589	27.931	28.009	29.229	28.759	28.405	27.945	25.912	26.619
##	1959	26.076	25.286	27.660	25.951	26.398	25.565	28.865	30.000	29.261	29.012	26.992	27.897

Comments

Note extras on `ts`:


- Time period is 1 year
- 12 observations per year (monthly) in `freq`
- First observation is 1st month of 1946 in `start`

Printing formats nicely.

Time plot

- Time plot shows extra pattern:

```
autoplot(ny.ts)
```




Comments on time plot

- steady increase (after initial drop)
- repeating pattern each year (seasonal component).
- Not stationary.

Differencing the New York births

Does differencing help here? Looks stationary, but some regular spikes:


```
ny.diff.ts=diff(ny.ts)  
autoplot(ny.diff.ts)
```



Decomposing a seasonal time series

A visual (using original data):

```
ny.d <- decompose(ny.ts)  
ny.d %>% autoplot()
```



Decomposition bits

Shows:

- original series
- a “seasonal” part: something that repeats every year
- just the trend, going steadily up (except at the start)
- random: what is left over (“remainder”)

The seasonal part

Fitted seasonal part is same every year, births lowest in February and highest in July:

```
ny.d$seasonal
```


##	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug
## 1946	-0.6771947	-2.0829607	0.8625232	-0.8016787	0.2516514	-0.1532556	1.4560457	1.1645938
## 1947	-0.6771947	-2.0829607	0.8625232	-0.8016787	0.2516514	-0.1532556	1.4560457	1.1645938
## 1948	-0.6771947	-2.0829607	0.8625232	-0.8016787	0.2516514	-0.1532556	1.4560457	1.1645938
## 1949	-0.6771947	-2.0829607	0.8625232	-0.8016787	0.2516514	-0.1532556	1.4560457	1.1645938
## 1950	-0.6771947	-2.0829607	0.8625232	-0.8016787	0.2516514	-0.1532556	1.4560457	1.1645938
## 1951	-0.6771947	-2.0829607	0.8625232	-0.8016787	0.2516514	-0.1532556	1.4560457	1.1645938
## 1952	-0.6771947	-2.0829607	0.8625232	-0.8016787	0.2516514	-0.1532556	1.4560457	1.1645938
## 1953	-0.6771947	-2.0829607	0.8625232	-0.8016787	0.2516514	-0.1532556	1.4560457	1.1645938
## 1954	-0.6771947	-2.0829607	0.8625232	-0.8016787	0.2516514	-0.1532556	1.4560457	1.1645938
## 1955	-0.6771947	-2.0829607	0.8625232	-0.8016787	0.2516514	-0.1532556	1.4560457	1.1645938
## 1956	-0.6771947	-2.0829607	0.8625232	-0.8016787	0.2516514	-0.1532556	1.4560457	1.1645938
## 1957	-0.6771947	-2.0829607	0.8625232	-0.8016787	0.2516514	-0.1532556	1.4560457	1.1645938
## 1958	-0.6771947	-2.0829607	0.8625232	-0.8016787	0.2516514	-0.1532556	1.4560457	1.1645938
## 1959	-0.6771947	-2.0829607	0.8625232	-0.8016787	0.2516514	-0.1532556	1.4560457	1.1645938
##	Sep	Oct	Nov	Dec				
## 1946	0.6916162	0.7752444	-1.1097652	-0.3768197				
## 1947	0.6916162	0.7752444	-1.1097652	-0.3768197				
## 1948	0.6916162	0.7752444	-1.1097652	-0.3768197				
## 1949	0.6916162	0.7752444	-1.1097652	-0.3768197				
## 1950	0.6916162	0.7752444	-1.1097652	-0.3768197				
## 1951	0.6916162	0.7752444	-1.1097652	-0.3768197				
## 1952	0.6916162	0.7752444	-1.1097652	-0.3768197				
## 1953	0.6916162	0.7752444	-1.1097652	-0.3768197				
## 1954	0.6916162	0.7752444	-1.1097652	-0.3768197				

Time Series

Time series basics: white noise

Each value independent random normal. Knowing one value tells you nothing about the next. “Random” process.

```
wn=rnorm(100)
wn.ts=ts(wn)
autoplot(wn.ts)
```



Lagging a time series

This means moving a time series one (or more) steps back in time:


```
x=rnorm(5)
tibble(x) %>% mutate(x_lagged=lag(x)) -> with_lagged
with_lagged
```

x	x_lagged
-2.0360948	NA
-0.5786158	-2.0360948
0.6083646	-0.5786158
0.1180334	0.6083646
0.0563443	0.1180334

Gain a missing because there is nothing before the first observation.

Lagging white noise

```
tibble(wn) %>% mutate(wn_lagged=lag(wn)) -> wn_with_lagged  
ggplot(wn_with_lagged, aes(y=wn, x=wn_lagged))+geom_point()
```



```
with(wn_with_lagged, cor.test(wn, wn_lagged, use="c")) # ignore
```

##

Correlation with lagged series

If you know about white noise at one time point, you know *nothing* about it at the next. This is shown by the scatterplot and the correlation.


On the other hand, this:

```
tibble(age=kings$X1) %>%  
  mutate(age_lagged=lag(age)) -> kings_with_lagged  
with(kings_with_lagged, cor.test(age, age_lagged))
```

```
##  
## Pearson's product-moment correlation  
##  
## data: age and age_lagged  
## t = 2.7336, df = 39, p-value = 0.00937  
## alternative hypothesis: true correlation is not equal to 0  
## 95 percent confidence interval:  
## 0.1064770 0.6308209  
## sample estimates:  
## cor  
## 0.4009919
```

Correlation with next value?

```
ggplot(kings_with_lagged, aes(x=age_lagged, y=age)) +  
  geom_point()
```



Two steps back:


```
kings_with_lagged %>%  
  mutate(age_lag_2=lag(age_lagged)) %>%  
  with(., cor.test(age, age_lag_2))  
  
##  
## Pearson's product-moment correlation  
##  
## data: age and age_lag_2  
## t = 1.5623, df = 38, p-value = 0.1265  
## alternative hypothesis: true correlation is not equal to 0  
## 95 percent confidence interval:  
## -0.07128917 0.51757510  
## sample estimates:  
## cor  
## 0.245676
```

Still a correlation two steps back, but smaller (and no longer significant).

Autocorrelation

Correlation of time series with *itself* one, two,... time steps back is useful idea, called **autocorrelation**. Make a plot of it with `acf` and `autoplot`. Here, white noise:

```
acf(wn.ts, plot=F) %>% autoplot()
```




No autocorrelations beyond chance, anywhere (except *possibly* at lag 13).

Autocorrelations work best on stationary series

Kings, differenced

```
acf(kings.diff.ts, plot=F) %>% autoplot()
```




Comments on autocorrelations of kings series

Negative autocorrelation at lag 1, nothing beyond that.

- If one value of differenced series positive, next one most likely negative.
- If one monarch lives longer than predecessor, next one likely lives shorter.

NY births, differenced

```
acf(ny.diff.ts, plot=F) %>% autoplot()
```



Lots of stuff:

- large positive autocorrelation at 1.0 years (July one year like July last year)
- large negative autocorrelation at 1 month.
- smallish but significant negative autocorrelation at 0.5 year = 6 months.
- Other stuff – complicated.

Souvenir sales


Monthly sales for a beach souvenir shop in Queensland, Australia:

```
souv=read_table("souvenir.txt", col_names=F)
souv.ts=ts(souv,frequency=12,start=1987)
souv.ts
```

##	Jan	Feb	Mar	Apr	May	Jun	Jul
## 1987	1664.81	2397.53	2840.71	3547.29	3752.96	3714.74	4349.61
## 1988	2499.81	5198.24	7225.14	4806.03	5900.88	4951.34	6179.12
## 1989	4717.02	5702.63	9957.58	5304.78	6492.43	6630.80	7349.62
## 1990	5921.10	5814.58	12421.25	6369.77	7609.12	7224.75	8121.22
## 1991	4826.64	6470.23	9638.77	8821.17	8722.37	10209.48	11276.55
## 1992	7615.03	9849.69	14558.40	11587.33	9332.56	13082.09	16732.78
## 1993	10243.24	11266.88	21826.84	17357.33	15997.79	18601.53	26155.15
##	Aug	Sep	Oct	Nov	Dec		
## 1987	3566.34	5021.82	6423.48	7600.60	19756.21		
## 1988	4752.15	5496.43	5835.10	12600.08	28541.72		
## 1989	8176.62	8573.17	9690.50	15151.84	34061.01		
## 1990	7979.25	8093.06	8476.70	17914.66	30114.41		
## 1991	12552.22	11637.39	13606.89	21822.11	45060.69		
## 1992	19888.61	23933.38	25391.35	36024.80	80721.71		
## 1993	28586.52	30505.41	30821.33	46634.38	104660.67		

Plot of souvenir sales

```
autoplot(souv.ts)
```




Several problems:

- Mean goes up over time
- Variability gets larger as mean gets larger
- Not stationary

Problem-fixing:


Fix non-constant variability first by taking logs:

```
souv.log.ts=log(souv.ts)  
autoplot(souv.log.ts)
```



Mean still not constant, so try taking differences

```
souv.log.diff.ts=diff(souv.log.ts)  
autoplot(souv.log.diff.ts)
```




Comments

- Now stationary
- but clear seasonal effect.

Decomposing to see the seasonal effect

```
souv.d=decompose(souv.log.diff.ts)  
autoplot(souv.d)
```



Comments

Big drop in one month's differences. Look at seasonal component to see which:


```
souv.d$seasonal
```

##	Jan	Feb	Mar	Apr	May	Jun	Jul
## 1987		0.23293343	0.49068755	-0.39700942	0.02410429	0.05074206	0.13552988
## 1988	-1.90372141	0.23293343	0.49068755	-0.39700942	0.02410429	0.05074206	0.13552988
## 1989	-1.90372141	0.23293343	0.49068755	-0.39700942	0.02410429	0.05074206	0.13552988
## 1990	-1.90372141	0.23293343	0.49068755	-0.39700942	0.02410429	0.05074206	0.13552988
## 1991	-1.90372141	0.23293343	0.49068755	-0.39700942	0.02410429	0.05074206	0.13552988
## 1992	-1.90372141	0.23293343	0.49068755	-0.39700942	0.02410429	0.05074206	0.13552988
## 1993	-1.90372141	0.23293343	0.49068755	-0.39700942	0.02410429	0.05074206	0.13552988
##	Aug	Sep	Oct	Nov	Dec		
## 1987	-0.03710275	0.08650584	0.09148236	0.47311204	0.75273614		
## 1988	-0.03710275	0.08650584	0.09148236	0.47311204	0.75273614		
## 1989	-0.03710275	0.08650584	0.09148236	0.47311204	0.75273614		
## 1990	-0.03710275	0.08650584	0.09148236	0.47311204	0.75273614		
## 1991	-0.03710275	0.08650584	0.09148236	0.47311204	0.75273614		
## 1992	-0.03710275	0.08650584	0.09148236	0.47311204	0.75273614		
## 1993	-0.03710275	0.08650584	0.09148236	0.47311204	0.75273614		

January.

Autocorrelations

```
acf(souv.log.diff.ts, plot=F) %>% autoplot()
```



- Big positive autocorrelation at 1 year (strong seasonal effect)
- Small negative autocorrelation at 1 and 2 months.

Moving average

- A particular type of time series called a **moving average** or MA process captures idea of autocorrelations at a few lags but not at others.
- Here's generation of MA(1) process, with autocorrelation at lag 1 but not otherwise:

```
beta=1
tibble(e=rnorm(100)) %>%
  mutate(e_lag=lag(e)) %>%
  mutate(y=e+beta*e_lag) %>%
  mutate(y=ifelse(is.na(y), 0, y)) -> ma
```

The series

ma

e	e_lag	y
0.0779806	NA	0.0000000
1.6644479	0.0779806	1.7424284
-1.4539253	1.6644479	0.2105226
-0.4015166	-1.4539253	-1.8554419
0.6809716	-0.4015166	0.2794549
0.4051565	0.6809716	1.0861281
0.2755077	0.4051565	0.6806642
-0.1823334	0.2755077	0.0931743
0.2264065	-0.1823334	0.0440731
0.3606240	0.2264065	0.5870305
2.2764053	0.3606240	2.6370293
-1.7780947	2.2764053	0.4983106
0.9387412	-1.7780947	-0.8393535
0.8939353	0.9387412	1.8326765
-0.1715134	0.8939353	0.7224219
0.7056000	-0.1715134	0.5641000

Time Series


Comments

- e contains independent “random shocks”.
- Start process at 0.
- Then, each value of the time series has that time's random shock, plus a multiple of the last time's random shock.
- $y[i]$ has shock in common with $y[i-1]$; should be a lag 1 autocorrelation.
- But $y[i]$ has no shock in common with $y[i-2]$, so no lag 2 autocorrelation (or beyond).

ACF for MA(1) process

Significant at lag 1, but beyond, just chance:

```
acf(ma$y, plot=F, na.rm=T) %>% autoplot()
```



AR process

Another kind of time series is AR process, where each value depends on previous one, like this (loop):

```
e=rnorm(100)
x=numeric(0)
x[1]=0
alpha=0.7
for (i in 2:100)
{
  x[i]=alpha*x[i-1]+e[i]
}
```


The series

x

##	[1]	0.00000000	0.48114799	0.71103525
##	[4]	0.65900463	0.90960457	1.12185864
##	[7]	1.02390502	-0.69991592	-0.64803277
##	[10]	0.23657047	0.16418958	-0.15043538
##	[13]	-0.09667020	0.88635091	1.65971423
##	[16]	2.62912167	1.43019873	1.45511765
##	[19]	3.11845464	2.95303071	0.36676230
##	[22]	0.22714525	0.48741737	-1.12005103
##	[25]	0.83144302	-0.07876862	0.52370605
##	[28]	-0.32393795	-0.31129337	2.06203136
##	[31]	1.74299095	1.93791340	1.04509322
##	[34]	0.68711330	1.83912870	1.06043342
##	[37]	2.56960344	1.72169161	1.23413651
##	[40]	1.17561493	3.12403601	1.58765927
##	[43]	0.13744074	-0.05372973	-0.44291441


Time Series

Comments

- Each random shock now only used for its own value of x
- but $x[i]$ also depends on previous value $x[i-1]$
- so correlated with previous value
- *but* $x[i]$ also contains multiple of $x[i-2]$ and previous x 's
- so all x 's correlated, but autocorrelation dying away.

ACF for AR(1) series


```
acf(x, plot=F) %>% autoplot()
```



Partial autocorrelation function

This cuts off for an AR series:


```
pacf(x, plot=F) %>% autoplot()
```



The lag-2 autocorrelation should not be significant, and isn't.

PACF for an MA series decays slowly

```
pacf(ma$y, plot=F) %>% autoplot()
```



The old way of doing time series analysis

Starting from a series with constant variability (eg. transform first to get it, as for souvenirs):

- Assess stationarity.
- If not stationary, take differences as many times as needed until it is.
- Look at ACF, see if it dies off. If it does, you have MA series.
- Look at PACF, see if that dies off. If it does, have AR series.
- If neither dies off, probably have a mixed “ARMA” series.
- Fit coefficients (like regression slopes).
- Do forecasts.

The new way of doing time series analysis (in R)

- Transform series if needed to get constant variability
- Use package `forecast`.
- Use function `auto.arima` to estimate what kind of series best fits data.
- Use `forecast` to see what will happen in future.

Anatomy of auto.arima output

```
auto.arima(ma$y)
```

```
## Series: ma$y
## ARIMA(5,0,0) with zero mean
##
## Coefficients:
##          ar1          ar2          ar3          ar4          ar5
##          0.8477   -0.7214   0.5633   -0.4953   0.1973
## s.e.    0.0988    0.1221   0.1289    0.1239   0.1037
##
## sigma^2 estimated as 1.273:  log likelihood=-152.06
## AIC=316.12   AICc=317.03   BIC=331.76
```

Comments over.

Comments

- ARIMA part tells you what kind of series you are estimated to have:
 - first number (first 0) is AR (autoregressive) part
 - second number (second 0) is amount of differencing here
 - third number (1) is MA (moving average) part
- Below that, coefficients (with SEs)
- AICc is measure of fit (lower better)

What other models were possible?

Run `auto.arima` with `trace=T`:

```
auto.arima(ma$y, trace=T)
```

```
##  
## ARIMA(2,0,2) with non-zero mean : Inf  
## ARIMA(0,0,0) with non-zero mean : 365.3271  
## ARIMA(1,0,0) with non-zero mean : 342.1337  
## ARIMA(0,0,1) with non-zero mean : Inf  
## ARIMA(0,0,0) with zero mean : 363.2452  
## ARIMA(2,0,0) with non-zero mean : 332.1106  
## ARIMA(3,0,0) with non-zero mean : 329.8212  
## ARIMA(4,0,0) with non-zero mean : 320.5205  
## ARIMA(5,0,0) with non-zero mean : 319.3257  
## ARIMA(5,0,1) with non-zero mean : Inf  
## ARIMA(4,0,1) with non-zero mean : Inf  
## ARIMA(5,0,0) with zero mean : 317.0272  
## ARIMA(4,0,0) with zero mean : 318.2995  
## ARIMA(5,0,1) with zero mean : Inf
```

Doing it all the new way: white noise

```
wn.aa=auto.arima(wn.ts)
wn.aa
```

```
## Series: wn.ts
## ARIMA(0,0,0) with zero mean
##
## sigma^2 estimated as 1.111:  log likelihood=-147.16
## AIC=296.32   AICc=296.36   BIC=298.93
```

Best fit *is* white noise (no AR, no MA, no differencing).

Forecasts:

```
forecast(wn.aa)
```

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## 101	0	-1.350869	1.350869	-2.065975	2.065975
## 102	0	-1.350869	1.350869	-2.065975	2.065975
## 103	0	-1.350869	1.350869	-2.065975	2.065975
## 104	0	-1.350869	1.350869	-2.065975	2.065975
## 105	0	-1.350869	1.350869	-2.065975	2.065975
## 106	0	-1.350869	1.350869	-2.065975	2.065975
## 107	0	-1.350869	1.350869	-2.065975	2.065975
## 108	0	-1.350869	1.350869	-2.065975	2.065975
## 109	0	-1.350869	1.350869	-2.065975	2.065975
## 110	0	-1.350869	1.350869	-2.065975	2.065975

Forecasts all 0, since the past doesn't help to predict future.

MA(1)

```
y.aa=auto.arima(ma$y)
y.aa
```


```
## Series: ma$y
## ARIMA(5,0,0) with zero mean
##
## Coefficients:
##          ar1          ar2          ar3          ar4          ar5
##          0.8477   -0.7214   0.5633   -0.4953   0.1973
## s.e.  0.0988    0.1221   0.1289    0.1239   0.1037
##
## sigma^2 estimated as 1.273:  log likelihood=-152.06
## AIC=316.12   AICc=317.03   BIC=331.76

y.f=forecast(y.aa)
```

Plotting the forecasts for MA(1)

```
autoplot(y.f)
```

Forecasts from ARIMA(5,0,0) with zero mean



AR(1)

```
x.aa=auto.arima(x)
x.aa
```

```
## Series: x
## ARIMA(1,0,0) with non-zero mean
##
## Coefficients:
##           ar1      mean
##          0.6843  0.4811
## s.e.  0.0720  0.2866
##
## sigma^2 estimated as 0.8717:  log likelihood=-134.33
## AIC=274.66   AICc=274.91   BIC=282.48
```

Oops! Thought it was MA(1), not AR(1)!

Fit right AR(1) model:

```
x.arima=arima(x,order=c(1,0,0))  
x.arima
```

```
##
```

```
## Call:
```

```
## arima(x = x, order = c(1, 0, 0))
```

```
##
```

```
## Coefficients:
```

```
##          ar1  intercept
```

```
##          0.6843      0.4811
```

```
## s.e.    0.0720      0.2866
```


```
##
```

```
## sigma^2 estimated as 0.8542:  log likelihood = -134.33,  ai
```


Forecasts for x

```
forecast(x.arima) %>% autoplot()
```


Forecasts from ARIMA(1,0,0) with non-zero mean



Comparing wrong model:

```
forecast(x.aa) %>% autoplot()
```

Forecasts from ARIMA(1,0,0) with non-zero mean



Kings

```
kings.aa=auto.arima(kings.ts)
kings.aa
```

```
## Series: kings.ts
## ARIMA(0,1,1)
##
## Coefficients:
##          ma1
##      -0.7218
## s.e.    0.1208
##
## sigma^2 estimated as 236.2:  log likelihood=-170.06
## AIC=344.13   AICc=344.44   BIC=347.56
```


Kings forecasts:

```
kings.f=forecast(kings.aa)
kings.f
```

##	Point	Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## 43		67.75063	48.05479	87.44646	37.62845	97.87281
## 44		67.75063	47.30662	88.19463	36.48422	99.01703
## 45		67.75063	46.58489	88.91637	35.38042	100.12084
## 46		67.75063	45.88696	89.61429	34.31304	101.18822
## 47		67.75063	45.21064	90.29062	33.27869	102.22257
## 48		67.75063	44.55402	90.94723	32.27448	103.22678
## 49		67.75063	43.91549	91.58577	31.29793	104.20333
## 50		67.75063	43.29362	92.20763	30.34687	105.15439
## 51		67.75063	42.68718	92.81408	29.41939	106.08187
## 52		67.75063	42.09507	93.40619	28.51383	106.98742

Kings forecasts, plotted

```
autoplot(kings.f) + labs(x="index", y= "age at death")
```



NY births

Very complicated:

```
ny.aa=auto.arima(ny.ts)
ny.aa
```

```
## Series: ny.ts
## ARIMA(2,1,2)(1,1,1)[12]
##
## Coefficients:
##          ar1      ar2      ma1      ma2      sar1      sma1
##          0.6539 -0.4540 -0.7255  0.2532 -0.2427 -0.8451
## s.e.  0.3003   0.2429   0.3227  0.2878   0.0985   0.0995
##
## sigma^2 estimated as 0.4076:  log likelihood=-157.45
## AIC=328.91   AICc=329.67   BIC=350.21
```

NY births forecasts

Not *quite* same every year:


```
ny.f=forecast(ny.aa,h=36)  
ny.f
```

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## Jan 1960	27.69056	26.87069	28.51043	26.43668	28.94444
## Feb 1960	26.07680	24.95838	27.19522	24.36632	27.78728
## Mar 1960	29.26544	28.01566	30.51523	27.35406	31.17683
## Apr 1960	27.59444	26.26555	28.92333	25.56208	29.62680
## May 1960	28.93193	27.52089	30.34298	26.77392	31.08995
## Jun 1960	28.55379	27.04381	30.06376	26.24448	30.86309
## Jul 1960	29.84713	28.23370	31.46056	27.37960	32.31466
## Aug 1960	29.45347	27.74562	31.16132	26.84155	32.06539
## Sep 1960	29.16388	27.37259	30.95517	26.42433	31.90342
## Oct 1960	29.21343	27.34498	31.08188	26.35588	32.07098
## Nov 1960	27.26221	25.31879	29.20563	24.29000	30.23441
## Dec 1960	28.06863	26.05137	30.08589	24.98349	31.15377
## Jan 1961	27.66908	25.59684	29.74132	24.49986	30.83830
## Feb 1961	26.21255	24.08615	28.33895	22.96051	29.46460
## Mar 1961	29.22612	27.04347	31.40878	25.88804	32.56420
## Apr 1961	27.58011	25.34076	29.81945	24.15533	31.00488
## May 1961	28.71354	26.41925	31.00783	25.20473	32.22235
## Jun 1961	28.21736	25.87042	30.56429	24.62803	31.80668
## Jul 1961	29.98728	27.58935	32.38521	26.31996	33.65460
## Aug 1961	29.96127	27.51330	32.40925	26.21743	33.70512
## Sep 1961	29.56515	27.06786	32.06243	25.74588	33.38441
## Oct 1961	29.54543	26.99965	32.09121	25.65200	33.43886
## Nov 1961	27.57845	24.98510	30.17181	23.61226	31.54465
## Dec 1961	28.40796	25.76792	31.04801	24.37036	32.44556
## Jan 1962	28.05431	25.33756	30.77106	23.88938	32.20922

Time Series

Plotting the forecasts

```
autoplot(ny.f)+labs(x="time", y="births")
```



Log-souvenir sales

```
souv.aa=auto.arima(souv.log.ts)
souv.aa
```

```
## Series: souv.log.ts
## ARIMA(2,0,0)(0,1,1)[12] with drift
##
## Coefficients:
##          ar1      ar2      sma1    drift
##      0.3470  0.3516  -0.5205  0.0238
## s.e.  0.1092  0.1115   0.1700  0.0031
##
## sigma^2 estimated as 0.02953:  log likelihood=24.54
## AIC=-39.09   AICc=-38.18   BIC=-27.71
```

```
souv.f=forecast(souv.aa,h=27)
```

The forecasts

Differenced series showed low value for January (large drop). December highest, Jan and Feb lowest:


```
souv.f
```

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## Jan 1994	9.578291	9.358036	9.798545	9.241440	9.915141
## Feb 1994	9.754836	9.521700	9.987972	9.398285	10.111386
## Mar 1994	10.286195	10.030937	10.541453	9.895811	10.676578
## Apr 1994	10.028630	9.765727	10.291532	9.626555	10.430704
## May 1994	9.950862	9.681555	10.220168	9.538993	10.362731
## Jun 1994	10.116930	9.844308	10.389551	9.699991	10.533868
## Jul 1994	10.369140	10.094251	10.644028	9.948734	10.789545
## Aug 1994	10.460050	10.183827	10.736274	10.037603	10.882498
## Sep 1994	10.535595	10.258513	10.812677	10.111835	10.959356
## Oct 1994	10.585995	10.308386	10.863604	10.161429	11.010561
## Nov 1994	11.017734	10.739793	11.295674	10.592660	11.442807
## Dec 1994	11.795964	11.517817	12.074111	11.370575	12.221353
## Jan 1995	9.840884	9.540241	10.141527	9.381090	10.300678
## Feb 1995	10.015540	9.711785	10.319295	9.550987	10.480093
## Mar 1995	10.555070	10.246346	10.863794	10.082918	11.027222
## Apr 1995	10.299676	9.989043	10.610309	9.824604	10.774749
## May 1995	10.225535	9.913326	10.537743	9.748053	10.703017

Plotting the forecasts

```
autoplot(souv.f)
```

Forecasts from ARIMA(2,0,0)(0,1,1)[12] with drift



Global mean temperatures, revisited

```
temp.ts=ts(temp$temperature,start=1880)
temp.aa=auto.arima(temp.ts)
temp.aa
```

```
## Series: temp.ts
## ARIMA(1,1,3) with drift
##
## Coefficients:
##          ar1      ma1      ma2      ma3      drift
##      -0.9374  0.5038  -0.6320  -0.2988  0.0067
## s.e.   0.0835  0.1088  0.0876  0.0844  0.0025
##
## sigma^2 estimated as 0.008939:  log likelihood=124.34
## AIC=-236.67   AICc=-235.99   BIC=-219.47
```

Forecasts

```
temp.f=forecast(temp.aa)  
autoplot(temp.f)+labs(x="year", y="temperature")
```

Forecasts from ARIMA(1,1,3) with drift

